

proj10: Virus Computational Model

Background: Viruses, Drug Treatments, and Computational Models

Viruses such as HIV and H1N1 represent a significant challenge to modern medicine. One of the reasons that they are so difficult to treat is because of their ability to evolve.

As you may know from science classes, the traits of an organism are determined by its genetic code. When organisms reproduce, their offspring will inherit genetic information from their parent. This genetic information will be modified, either because of mixing of the two parents' genetic information, or through mutations in the genome replication process, thus introducing diversity into a population.

Viruses are no exception. Two characteristics of viruses make them particularly difficult to treat.

The first is that their replication mechanism often lacks the error checking mechanisms that are present in more complex organisms. This speeds up the rate of mutation. Secondly, viruses replicate extremely quickly, orders of magnitude faster than humans. Thus, while we may be used to thinking of evolution as a process which occurs over long time scales, populations of viruses can undergo substantial evolutionary changes within a single patient over the course of treatment.

These two characteristics allow a virus population to quickly acquire genetic resistance to therapy. In this problem set, we will make use of simulations to explore the effect of introducing drugs on the virus population and determine how best to address these treatment challenges within a simplified model.

Computational modeling has played an important role in the study of viruses such as HIV. In this project, we will implement a highly simplified stochastic model of virus population dynamics. Many details have been swept under the rug (host cells are not explicitly modeled and the size of the population is several orders of magnitude less than the size of actual virus populations). Nevertheless, our model exhibits biologically relevant characteristics and will give you a chance to analyze and interpret interesting simulation data.

Spread of a Virus in a Person

In reality, diseases are caused by viruses and have to be treated with medicine, so in the remainder of this problem set, we'll be looking at a detailed simulation of the spread of a virus within a Person. We've provided you with skeleton code in `proj10.py`.

Problem 1: Implementing a Simple Simulation (No Drug Treatment)

We start with a trivial model of the virus population – the patient does not take any drugs and the viruses do not acquire resistance to drugs. We simply model the virus population inside a patient as if it were left untreated.

To implement this model, you will need to fill in the SimpleVirus class, which maintains the state of a single virus particle. You will implement the methods `__init__`, `doesClear`, and `reproduce` according to the specifications. Use `random.random()` for generating random numbers to ensure that your results are consistent with ours.

Hint: during debugging, you might want to use `random.seed()` so that your results are reproducible.

The `reproduce` method in SimpleVirus should produce an offspring by returning a new instance of SimpleVirus with probability:

```
self.maxBirthProb * (1 – popDensity)
```

`self.maxBirthProb` is the birth rate under optimal conditions (the virus population is negligible relative to the available host cells so there is ample nourishment available). `popDensity` is defined as the ratio of the current virus population to the maximum virus population for a patient and should be calculated in the `update` method of the SimplePatient class.

You will also need to implement the SimplePatient class, which maintains the state of a virus population associated with a patient.

The `update` method in the SimplePatient class is the inner loop of the simulation. It modifies the state of the virus population for a single time step and returns the total virus population at the end of the time step. At every time step of the simulation, each virus particle has a fixed probability of being cleared (eliminated from the patient's body). Then if the virus particle is not cleared, it is considered for reproduction.

(Unlike the clearance probability, which is constant, the probability of a virus particle reproducing is a function of the virus population. With a larger virus population, there are fewer resources in the patient's body to facilitate reproduction, and the probability of reproduction will be lower. One way to think of this limitation is to consider that virus particles need to make use of a patient's cells to reproduce, they cannot reproduce on their own. As the virus population increases, there will be fewer available host cells for viruses to utilize for reproduction.)

To summarize, `update` should first decide which virus particles are cleared and which survive by making use of the `doesClear` method of each SimpleVirus instance and update the collection of SimpleVirus instances accordingly. With the surviving SimpleVirus instances, `update` should then call the `reproduce` method for each virus particle. Based on the population density of the surviving SimpleVirus instances, `reproduce` should either return a new instance of SimpleVirus representing the offspring of the virus particle, or raise a `NoChildException` indicating that the virus particle does not reproduce during the

current time step. The update method should update the attributes of the patient appropriately under either of these conditions.

After iterating through all the virus particles, the update method returns the number of virus particles in the patient at the end of the time step.

HINT: Be very wary about mutating an object while iterating over its elements. It is best to avoid this entirely (consider introducing additional “helper” variables).

Note that the mapping between time steps and actual time will vary depending on the type of virus being considered, but for this project, think of a time step as a simulated hour of time.

You will test your implementation in problem 2.

Problem 2: Running and Analyzing a Simple Simulation (No Drug Treatment)

You should start by understanding the population dynamics before introducing any drug. Fill in the function `simulationWithoutDrug()`. This method should instantiate a `SimplePatient` and repeatedly call the update method to simulate changes in the virus population over time. Save the population values over the course of the simulation and use `pylab` to plot the virus population as a function of time. Be sure to title and label your plot.

`SimplePatient` should be instantiated with the following parameters:

- `viruses`, a list of 100 `SimpleVirus` instances
- `maxPop`, Maximum Sustainable Virus Population = 1000

Each `SimpleVirus` instance in the `viruses` list should be initialized with the following parameters:

- `maxBirthProb`, Maximum Reproduction Probability for a Virus Particle = 0.1
- `clearProb`, Maximum Clearance Probability for a Virus Particle = 0.05

Fill in the function `simulationWithoutDrug()` that instantiates a patient, simulates changes to the virus population for 300 time steps (i.e. 300 calls to `update`), and plots the virus population as a function of time; that is, the x-axis should correspond the number of elapsed time steps, and y-axis the population of the virus in the patient. Run the simulation many times (without setting the seed). Produce a single plot that is representative of the average case as a result of running the simulation many times. Don't forget to include axes labels, a key for the curves, and a title on your plot.

Reference materials for `pylab`:

https://matplotlib.org/users/pyplot_tutorial.html

About how long does it take before the population stops growing?

HINT: Compared to the previous problem sets, testing your simulation code is more challenging, because the behavior of the code is stochastic, and the expected output is not exactly known. How do you know whether your plots are correct or not? One way to test this is to run the simulation with extreme input values (i.e. initialization parameters), and check that the output matches your intuition. For example, if `codeMaxBirthProb` is set to 0.99 instead of 0.1, then you would expect that the virus population rapidly increases over a short period of time. Similarly, if you run your simulation with `clearProb` = 0.99 and `maxBirthProb` = 0.1, then you should see the virus population quickly decreasing within a small number of steps. You can also try to vary the input values, and check whether the output plots change as you expect. For example, if you run multiple simulation runs, each time increasing `maxBirthProb`, the curves in the successive plots should show an “upward” trend, since the virus will reproduce faster with a higher `maxBirthProb`.